

Testing and Quality Analysis of commons-imaging Project

GitHub Repository

Nima Mehranfar

January 7, 2025

Abstract

This report presents the testing and quality analysis conducted on the commons-imaging project hosted on GitHub. The project undergoes thorough software quality analysis, CI/CD pipeline setup, code coverage analysis, performance testing, mutation testing, security assessment, and more. The results of these analyses are presented along with the actions taken to improve the quality of the commons-imaging codebase.

Contents

1	Introduction	4
1.1	Objectives	4
2	CI/CD Pipeline	5
2.1	CI/CD Configuration	5
3	Software Quality Analysis with SonarCloud	6
3.1	Issues Categorization	6
3.2	Refactoring and Rationale	6
4	Docker Image Creation and Containerization	12
4.1	Docker Image and Containerization	12
4.2	Dockerfile for Frontend	12
4.3	Dockerfile for Backend	12
4.4	Docker Compose Configuration	13
4.5	DockerHub Synchronization	13
4.5.1	Pulling Docker Images	14
4.5.2	Running Docker Images	14
4.5.3	Accessing the Web Application	14
5	Code Coverage Analysis	15
5.1	Setting Up JaCoCo	15
5.2	Initial Coverage Results	15
5.3	Final Coverage Results After Test Improvements	16
6	Mutation Testing with PiTest	17
6.1	Mutation Test Results	17
6.2	Detailed Analysis: ColorTools.java	18
7	Performance Testing with JMH	19
7.1	Benchmark Setup	19
7.2	Stress Testing Results	20
8	Automated Test Generation	21
8.1	Low Coverage Packages	21
8.2	Test Generation Process	21
8.3	Running Randoop	21
8.4	Test Coverage Improvement	22
8.5	Challenges and Limitations	23
9	Security Analysis with Snyk	24
9.1	Steps to Run Snyk	24
9.2	Vulnerability Report	24
9.3	Security Issues Found	24
10	Conclusion	27
11	References	28

List of Figures

3.1	SonarCloud Issues Overview: 2325 issues detected, including blocker and high severity issues.	7
3.2	Example of missing assertions in test cases.	7
3.3	Example of confusing or duplicate method/variable names.	8
3.4	Example of missing break in switch case.	8
3.5	9 issues fixed	9
3.6	Build halted due to failed assertions and floating-point precision issue of this test case.	9
3.7	Unfixable naming issue	10
3.8	Unfixable naming issue tagged as "accepted"	11
4.1	Docker Architecture for the commons-imaging-webapp.	12
4.2	Commons-Imaging Test Web Application Interface.	14
5.1	Visualization of Initial Code Coverage Results.	15
5.2	Visualization of Final Code Coverage Results.	16
6.1	Visualization of PIT Coverage Results.	17
6.2	Initial Coverage and Mutation Results for <code>ColorTools.java</code>	18
6.3	Improved Coverage and Mutation Results for <code>ColorTools.java</code>	18
7.1	Benchmark Results of Selected Methods of Commons-Imaging Project. .	20
8.1	JaCoCo report showing low coverage in specific packages.	21
8.2	Coverage improvement after generating tests using Randoop.	22
9.1	Snyk report showing no vulnerabilities for commons-imaging.	24
9.2	Vulnerabilities reduced after updating the OpenJDK version.	24
9.3	Unsanitized input vulnerability found in <code>SurveyTiffFolder.java:collectPaths</code> . .	25

List of Tables

5.1	Initial Code Coverage Results.	15
5.2	Final Code Coverage Results After Test Improvements.	16
6.1	Mutation Test Results Overview.	17
6.2	Mutation Test Results for <code>ColorTools.java</code>	18
7.1	Stress Testing Results of Commons-Imaging Project.	20
8.1	Coverage for <code>org.apache.commons.imaging.internal</code>	22
8.2	Coverage for <code>org.apache.commons.imaging.formats.psd.dataparsers</code>	22
8.3	Coverage for <code>org.apache.commons.imaging.formats.tiff</code>	22

1 Introduction

The *commons-imaging* project is a Java library developed by the Apache Software Foundation. It provides a comprehensive API for working with various image formats, enabling developers to read, write, and manipulate images with ease. This project focuses on evaluating the *commons-imaging* codebase through a rigorous quality assurance process to ensure reliability, maintainability, and performance.

1.1 Objectives

This testing project aims to achieve the following objectives:

- Establish CI/CD pipelines for the *commons-imaging* project using GitHub Actions.
- Perform software quality analysis leveraging SonarCloud for static code evaluation and reporting.
- Develop a Docker image to facilitate streamlined deployment and testing environments.
- Conduct code coverage analysis using JaCoCo to identify untested areas in the codebase.
- Apply mutation testing with PiTest to evaluate the robustness of existing test cases.
- Execute stress testing on project components with JMH to assess performance under load.
- Generate automated test cases for inadequately tested code segments using Randop.
- Perform security analysis using Snyk to detect and mitigate potential vulnerabilities.

2 CI/CD Pipeline

The *commons-imaging* project has been integrated into a CI/CD pipeline that can be built both locally and on cloud servers, including GitHub. This pipeline automates the build, test, and deployment processes to ensure continuous integration and delivery, enabling efficient and reliable software updates.

2.1 CI/CD Configuration

The CI/CD pipeline configuration utilizes GitHub Actions to automate the workflow. Below is the YAML file for the GitHub Actions workflow configuration:

Listing 2.1: GitHub Actions Workflow Configuration (`maven.yml`)

```
name: Java CI

on: [push, pull_request, workflow_dispatch]

permissions:
  contents: read

jobs:
  build:
    runs-on: ${{ matrix.os }}
    continue-on-error: ${{ matrix.experimental }}
    strategy:
      matrix:
        os: [windows-latest, macos-13, ubuntu-latest]
        java: [8, 11, 17, 21]
        experimental: [false]
      include:
        - java: 23
          os: ubuntu-latest
          experimental: false
        - java: 24-ea
          os: ubuntu-latest
          experimental: true

    steps:
      # Configure Git settings to avoid line-ending issues
      - name: Prepare git
        run: |
          git config --global core.autocrlf false
          git config --global core.eol lf
      - uses: actions/checkout@v2 # Checkout the repository code
        with:
          persist-credentials: false
      - name: Set up JDK ${{ matrix.java }}
        uses: actions/setup-java@v2 # Setup JDK version
        with:
          distribution: 'temurin'
          java-version: ${{ matrix.java }}
      - name: Build with Maven
        run: mvn --errors --show-version --batch-mode --no-transfer-progress
```

3 Software Quality Analysis with Sonar-Cloud

The software quality analysis of the *commons-imaging* project is conducted using Sonar-Cloud, a powerful tool that identifies potential code quality issues and provides recommendations for improvements. By integrating SonarCloud into the development workflow, developers can continuously monitor and address various code quality aspects to ensure a robust and maintainable codebase.

3.1 Issues Categorization

SonarCloud categorizes the issues it identifies into three main types:

- **Maintainability:** These refer to unnecessary complexity or suboptimal code patterns that may not be bugs but could lead to maintenance challenges.
- **Reliability:** Errors in the code that could lead to functional failures or unexpected behavior.
- **Security:** Potential security risks within the code that could be exploited by malicious actors.

3.2 Refactoring and Rationale

Several refactoring actions were undertaken to address the issues identified by Sonar-Cloud. For those issues that were skipped, a clear rationale was provided.

As illustrated in the image below (Figure 3.1), a total of 2325 issues were found, including 18 critical "blocker" severity issues and 333 "high" severity issues. All the "blocker" issues were maintainability; And I prioritized addressing some of the these issues to improve the overall code quality.

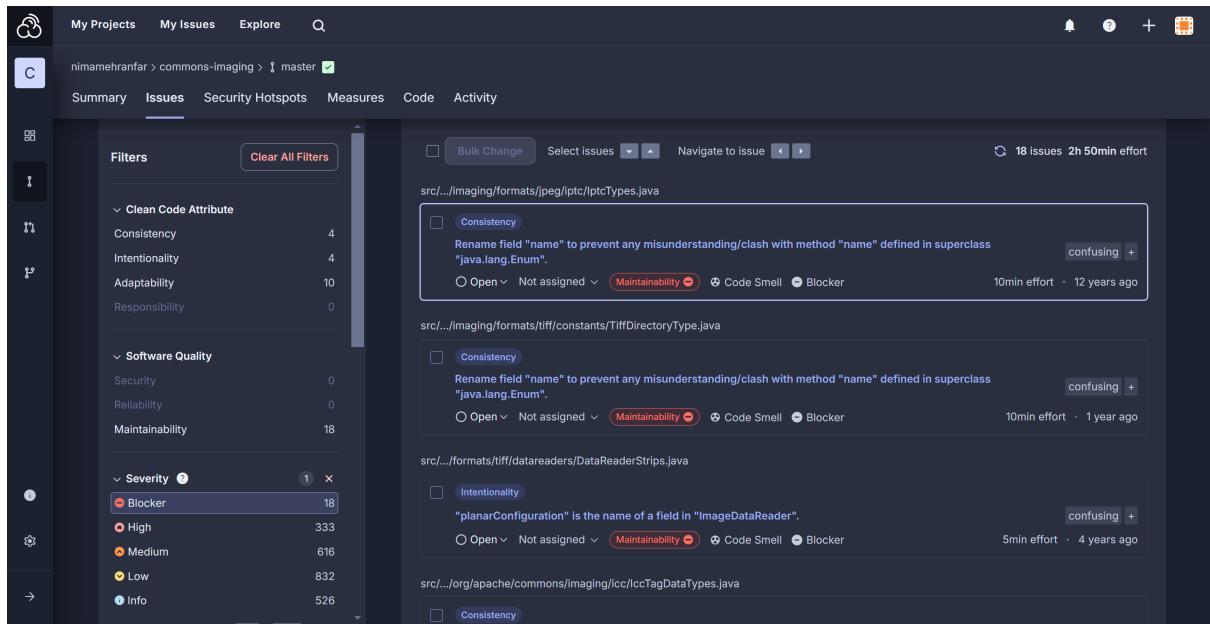


Figure 3.1: SonarCloud Issues Overview: 2325 issues detected, including blocker and high severity issues.

Some of the common issues included a lack of assertions in test cases, confusing or duplicate method/variable names and missing break in switch cases, as demonstrated in the images below (Figures 3.2, 3.3, 3.4).

The screenshot shows a Java code editor with four tabs: 'DataReaderStrips.java', 'RationalNumberTest.java' (which is the active tab), 'ImageDataReader.java', and 'DataReaderTiled.java'. The code in 'RationalNumberTest.java' is as follows:

```

33  public class RationalNumberTest extends AbstractImagingTest { ▲ Benedikt Ritter +7
35      public static Stream<Double> data() { ▲ Benedikt Ritter +4
...
105     -(Long.MAX_VALUE - 0.1) // ...
106     ).stream();
107 }
108
109     @ParameterizedTest ▲ Benedikt Ritter +4
110     @MethodSource("data")
111     public void testRationalNumber(final double testValue) {
112         final RationalNumber rational = RationalNumber.valueOf(testValue);
113         final double difference = Math.abs(testValue - rational.doubleValue());
114
115         final NumberFormat nf = NumberFormat.getInstance();
116         nf.setMaximumFractionDigits(15);
117
118         // TODO assert something here, the following will fail for some values. Do we have a bug?
119         // assertEquals(0.0, difference, 0.0);
120         Debug.debug("value: " + nf.format(testValue));
121         Debug.debug("rational: " + rational);
122         Debug.debug("difference: " + difference);
123         Debug.debug();
124     }
125
126     @Test ▲ gwlucastrig +1
127     public void testSpecialRationalNumber() {

```

The code editor interface includes line numbers on the left, a status bar at the bottom right showing '6 | 1 8', and a vertical scrollbar on the right.

Figure 3.2: Example of missing assertions in test cases.

```
19     public enum IptcTypes implements IptcType { 20 usages ± Gary Gregory +3
39         public static IptcType getUnknown(final int type) { 1 usage ± Damjan Jovanovic +3
40             return new IptcType() { ± Damjan Jovanovic +3
52     ↗     >         public String toString() { return "Unknown (" + type + ")"; }
55     >     }
56     }
57
58     public final int type;
59
60     public final String name;
61
62     IptcTypes(final int type, final String name) { 114 usages ± Gary Gregory
63         this.type = type;
64         this.name = name;
65     }
66
67     @Override ± Damjan Jovanovic +1
68     ↗     >     public String getName() {
69         return name;
70     }
71
72     @Override ± Damjan Jovanovic
73     ↗     >     public int getType() { return type; }
74
75     @Override ± Damjan Jovanovic +1
76     ↗     >     public String toString() { return name + " (" + type + ")"; }
77
78     ↗     >     }
79
80 }
```

Figure 3.3: Example of confusing or duplicate method/variable names.

```
44     public class ReadTagsAndImages { ± Gary Gregory +2
45
46     @     private static String interpretElements(final GeoKey geoKey, final int ref, final int len, final int valueOrPosition, fi
47     final String asciiParameters) {
48         switch (geoKey) {
49             case GTModelTypeGeoKey:
50                 switch (valueOrPosition) {
51                     case 1:
52                         return "Projected Coordinate System";
53                     case 2:
54                         return "Geographic Coordinate System";
55                     case 3:
56                         // the Geocentric coordinate system is seldom used
57                         return "Geocentric Coordinate System";
58                     default:
59                         break;
60                 }
61             case GTRasterTypeGeoKey:
62                 switch (valueOrPosition) {
63                     case 1:
64                         return "RasterPixelIsArea";
65                     case 2:
66                         return "RasterPixelIsPoint";
67                     default:
68                         return "User Defined";
69                 }
70         }
71     }
72 }
```

Figure 3.4: Example of missing break in switch case.

As shown in Figure 3.5, I successfully resolved 9 issues, but during the process, I also

introduced 1 additional issue which is related to commented-out lines of code in the failing test case assertion. This occurred due to a test case that revealed a bug in the code (as noted by the developer). The addition of assertions caused the build process to halt, which is evident in the error log below.

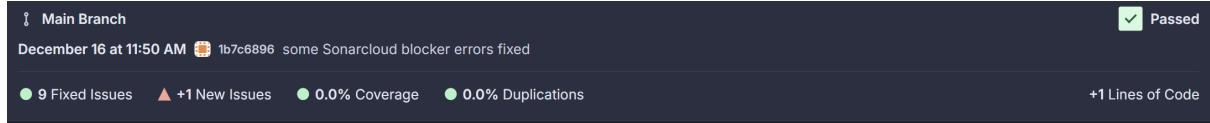


Figure 3.5: 9 issues fixed

The error message below demonstrates the issue with floating-point precision during the test execution:

```
Assertion Error: \\
[ERROR] Failures: \\
[ERROR] RationalNumberTest.testRationalNumber:120 The difference between the test
    ↪ value and the rational number representation exceeds the tolerance. \\
    ==> expected: <0.0> but was: <0.002000040531158447> \\
[INFO] \\
[ERROR] Tests run: 1076, Failures: 15, Errors: 0, Skipped: 7 \\
[INFO]

Patterns in Failures: \\
Small differences (<1e-9) might be acceptable due to floating-point arithmetic. \\
Large differences (e.g., 0.09999990463256836 or 9.223372034707292E18) are likely bugs
    ↪ or limitations in RationalNumber handling. \\
Therefore, the implementation of \texttt{rationalNumber.valueOf()} must be reviewed
    ↪ by the developer. \\
This problem was tagged as "Accepted" on SonarQube.
```

The screenshot shows an IDE editor with the file 'RationalNumberTest.java' open. The code is as follows:

```
33  public class RationalNumberTest extends AbstractImagingTest { ▲ Benedikt Ritter +7 *
35      public static Stream<Double> data() { ▲ Benedikt Ritter +4
36          return Stream.of(
37              -(Long.MAX_VALUE - 0.1) // ←
38          ).stream();
39      }
40
41      @ParameterizedTest ▲ Benedikt Ritter +4 *
42      @MethodSource("data")
43      public void testRationalNumber(final double testValue) {
44          final RationalNumber rational = RationalNumber.valueOf(testValue);
45          final double difference = Math.abs(testValue - rational.doubleValue());
46
47          final NumberFormat nf = NumberFormat.getInstance();
48          nf.setMaximumFractionDigits(15);
49
50          // Assert that the difference is within an acceptable tolerance
51          final double tolerance = 1e-15; // Define a small tolerance for floating-point comparison
52          assertEquals(expected: 0.0, difference, tolerance,
53              message: "The difference between the test value and the rational number representation exceeds the tolerance.");
54
55          Debug.debug("value: " + nf.format(testValue));
56          Debug.debug("rational: " + rational);
57          Debug.debug("difference: " + difference);
58          Debug.debug();
59      }
60
61      @Test
62      public void testRationalNumberWithLargeValue() {
63          final RationalNumber rational = RationalNumber.valueOf(9.223372034707292E18);
64          final double difference = Math.abs(9.223372034707292E18 - rational.doubleValue());
65
66          final NumberFormat nf = NumberFormat.getInstance();
67          nf.setMaximumFractionDigits(15);
68
69          // Assert that the difference is within an acceptable tolerance
70          final double tolerance = 1e-15; // Define a small tolerance for floating-point comparison
71          assertEquals(expected: 0.0, difference, tolerance,
72              message: "The difference between the test value and the rational number representation exceeds the tolerance.");
73
74          Debug.debug("value: " + nf.format(9.223372034707292E18));
75          Debug.debug("rational: " + rational);
76          Debug.debug("difference: " + difference);
77          Debug.debug();
78      }
79
80      @Test
81      public void testRationalNumberWithSmallValue() {
82          final RationalNumber rational = RationalNumber.valueOf(0.002000040531158447);
83          final double difference = Math.abs(0.002000040531158447 - rational.doubleValue());
84
85          final NumberFormat nf = NumberFormat.getInstance();
86          nf.setMaximumFractionDigits(15);
87
88          // Assert that the difference is within an acceptable tolerance
89          final double tolerance = 1e-15; // Define a small tolerance for floating-point comparison
90          assertEquals(expected: 0.0, difference, tolerance,
91              message: "The difference between the test value and the rational number representation exceeds the tolerance.");
92
93          Debug.debug("value: " + nf.format(0.002000040531158447));
94          Debug.debug("rational: " + rational);
95          Debug.debug("difference: " + difference);
96          Debug.debug();
97      }
98
99      @Test
100     public void testRationalNumberWithLargeValueWithLargeTolerance() {
101         final RationalNumber rational = RationalNumber.valueOf(9.223372034707292E18);
102         final double difference = Math.abs(9.223372034707292E18 - rational.doubleValue());
103
104         final NumberFormat nf = NumberFormat.getInstance();
105         nf.setMaximumFractionDigits(15);
106
107         // Assert that the difference is within an acceptable tolerance
108         final double tolerance = 1e-10; // Define a large tolerance for floating-point comparison
109         assertEquals(expected: 0.0, difference, tolerance,
110             message: "The difference between the test value and the rational number representation exceeds the tolerance.");
111
112         Debug.debug("value: " + nf.format(9.223372034707292E18));
113         Debug.debug("rational: " + rational);
114         Debug.debug("difference: " + difference);
115         Debug.debug();
116     }
117 }
```

The line 111 is highlighted with a red arrow pointing to it, indicating a failure. The status bar at the top right shows 3 errors and 5 warnings.

Figure 3.6: Build halted due to failed assertions and floating-point precision issue of this test case.

The issues related to floating-point precision need careful attention, as small discrepancies in numerical computations are often acceptable, but larger discrepancies indicate potential bugs. The developer has flagged the issue for further investigation; So I tagged this issue as "Accepted" on SonarQube.

Also in (Figures 3.7, 3.8), there is an issue labeled "child class fields should not shadow parent class fields" which is unfixable because the parent class field is a protected variable and cannot be accessed by the child class so any change to both the naming and/or syntax/keywords used for variables could result in an unwanted scenario and/or a bug. So I tagged this issue as "accepted" too.

A screenshot of a Java code editor showing a tooltip for a naming issue. The code is from `DataReaderStrips.java`. The tooltip points out that the field `planarConfiguration` is shadowing a field from the parent class `ImageDataReader`. It provides a link to the SonarQube rule description and options to fix or ignore the issue.

```
42 * See f@link ImageDataReader} for notes discussing design and development with particular emphasis 10 | 17 9 11 ^ v
43 */
44 public final class DataReaderStrips extends ImageDataReader { 5 usages  gwlucastrig +6
45
46     private final int bitsPerPixel; 12 usages
47     private final int compression; 7 usages
48     private final int rowsPerStrip; 16 usages
49     private final TiffPlanarConfiguration planarConfiguration; 2 usages
50     private final ByteOrder byteOrder; 5 usages
51     private int x; 8 usages
52     private int y; 19 usages
53     private final AbstractTiffImageData.Strips imageData
54
55     public DataReaderStrips(final TiffDirectory direc
56         final int[] bitsPerSample, final int pred
57         final int compression, final TiffPlanarConfiguration planarConfiguration, final ByteOrder byteOrder, final int row
58         final AbstractTiffImageData.Strips imageData) {
59         super(directory, photometricInterpreter, bitsPerSample, predictor, samplesPerPixel, sampleFormat, width, height, plan
60
61         this.bitsPerPixel = bitsPerPixel;
62         this.compression = compression;
63         this.rowsPerStrip = rowsPerStrip;
64         this.planarConfiguration = planarConfiguration;
65         this.imageData = imageData;
66         this.byteOrder = byteOrder;
```

"planarConfiguration" is the name of a field in "ImageDataReader".
SonarQube: Show rule description 'java:S2387' Alt+Shift+Enter More actions... Alt+Enter

org.apache.commons.imaging.formats.tiff.datareaders.DataReaderStrips
private final TiffPlanarConfiguration planarConfiguration
commons-imaging

Figure 3.7: Unfixable naming issue

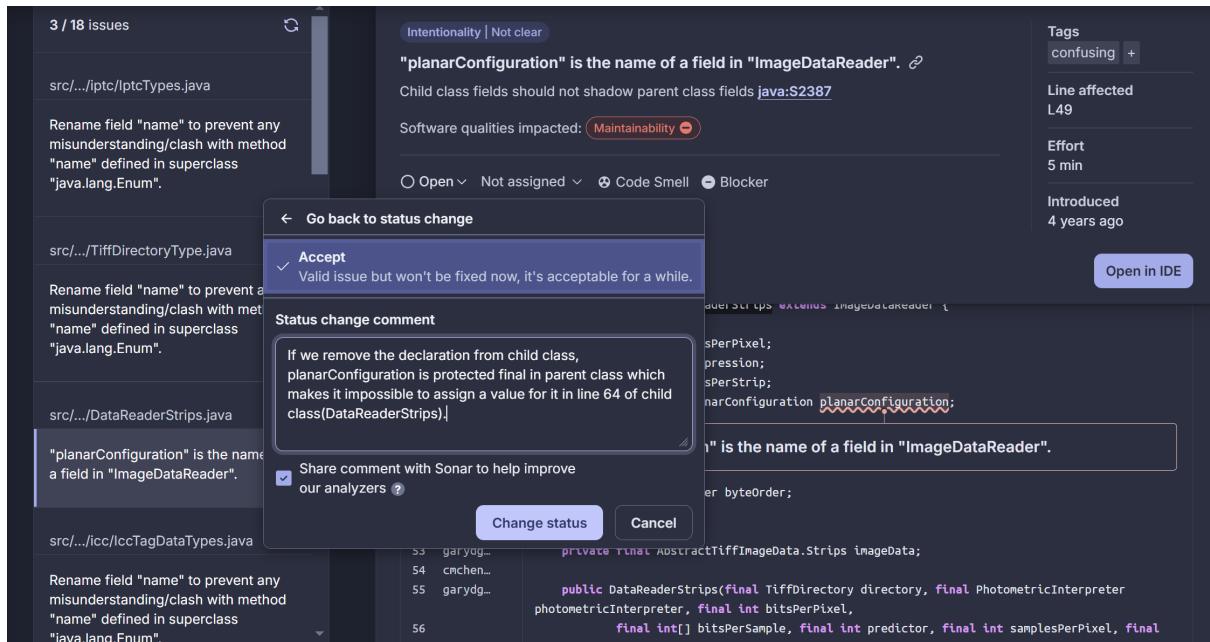


Figure 3.8: Unfixable naming issue tagged as "accepted"

4 Docker Image Creation and Containerization

The *commons-imaging* project, being an API and utility without a main method, required an additional web application for demonstration purposes. A simple web application was created using React.js for the frontend and Java Spring Boot for the backend. This setup allows the utility to be integrated into a functional system for testing and deployment.

4.1 Docker Image and Containerization

Two Docker images were created for the web application's frontend and backend. These images were pushed to DockerHub and configured to run together using Docker Compose. This ensures seamless orchestration and ease of deployment.

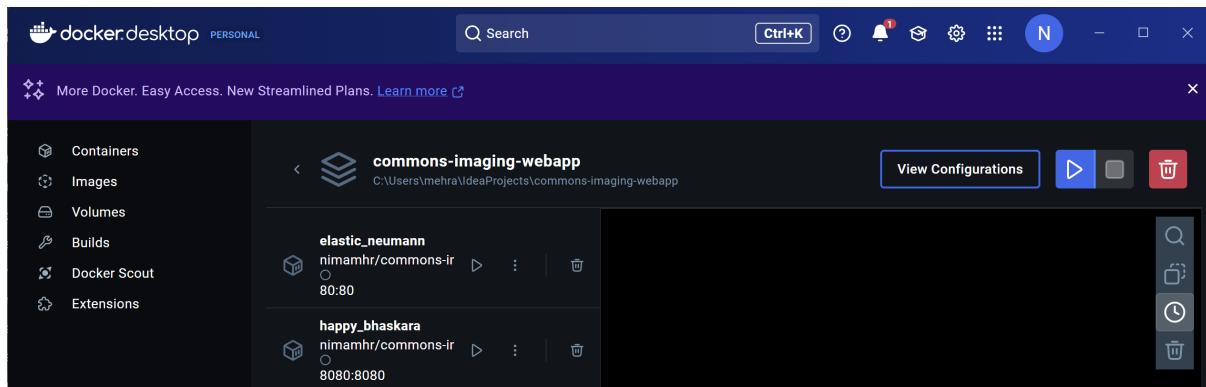


Figure 4.1: Docker Architecture for the commons-imaging-webapp.

4.2 Dockerfile for Frontend

The Dockerfile for the frontend builds a React.js application and serves it using Nginx. Below is the configuration:

Listing 4.1: Dockerfile for commons-imaging-webapp Frontend

```
FROM node:18-alpine AS build
WORKDIR /app
COPY package.json package-lock.json ./
RUN npm install
COPY . ./.
RUN npm run build
FROM nginx:alpine
COPY --from=build /app/build /usr/share/nginx/html
EXPOSE 80
CMD ["nginx", "-g", "daemon off;"]
```

4.3 Dockerfile for Backend

The backend is a Java Spring Boot application. The Dockerfile compiles the application and creates an executable JAR file:

Listing 4.2: Dockerfile for commons-imaging-webapp Backend

```
FROM openjdk:24-jdk
WORKDIR /app
COPY .mvn/.mvn
COPY mvnw pom.xml ./
RUN ./mvnw dependency:go-offline -B
COPY src/.src
RUN ./mvnw clean package -DskipTests
EXPOSE 8080
CMD ["java", "-jar", "target/test-0.0.1-SNAPSHOT.jar"]
```

4.4 Docker Compose Configuration

The Docker Compose file orchestrates both the frontend and backend services, setting up a shared network and dependencies.

Listing 4.3: Docker Compose File for Orchestration

```
version: '3.7'

services:
  backend:
    build:
      context: ./Backend
    ports:
      - "8080:8080"
    networks:
      - app-network

  frontend:
    build:
      context: ./commons-imaging-test
    ports:
      - "80:80"
    environment:
      - REACT_APP_API_URL=http://backend:8080 # API URL for backend
    networks:
      - app-network
    depends_on:
      - backend

networks:
  app-network:
    driver: bridge
```

4.5 DockerHub Synchronization

The Docker images for both the frontend and backend are synchronized with DockerHub for easy access. They can be pulled and run using the following commands:

4.5.1 Pulling Docker Images

Listing 4.4: Pull Docker Images from DockerHub

```
docker pull nimamhr/commons-imaging-webapp-frontend:latest  
docker pull nimamhr/commons-imaging-webapp-backend:latest
```

4.5.2 Running Docker Images

Listing 4.5: Run Docker Images

```
docker run -d -p 8080:8080 nimamhr/commons-imaging-webapp-backend:latest  
docker run -d -p 80:80 nimamhr/commons-imaging-webapp-frontend:latest
```

4.5.3 Accessing the Web Application

Once the containers are running, the web application can be accessed at:
<http://localhost/>

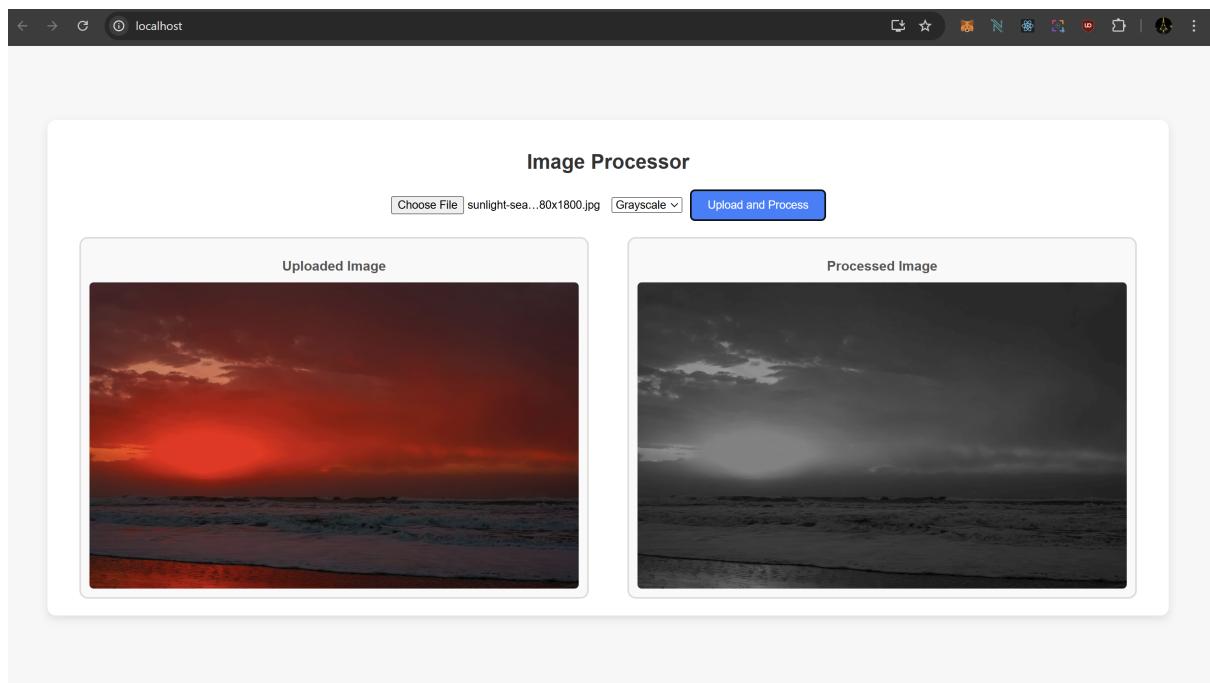


Figure 4.2: Commons-Imaging Test Web Application Interface.

5 Code Coverage Analysis

Code coverage analysis was performed using JaCoCo to assess the extent of test coverage within the *commons-imaging* project. This analysis helped identify untested or poorly tested portions of the codebase, guiding the addition of manual and automated tests.

5.1 Setting Up JaCoCo

To set up JaCoCo for the project, the JaCoCo Maven plugin was added to the `pom.xml` file. The following commands were then executed to generate code coverage reports:

Listing 5.1: Commands to Generate Code Coverage Reports

```
mvn clean test
mvn clean verify
```

5.2 Initial Coverage Results

The initial code coverage results, before adding tests, are summarized below:

Element	Missed Instructions	I. Cov. (%)	Missed Branches	B. Cov. (%)
Total	21,380 of 95,452	77%	2,570 of 7,199	64%

Table 5.1: Initial Code Coverage Results.

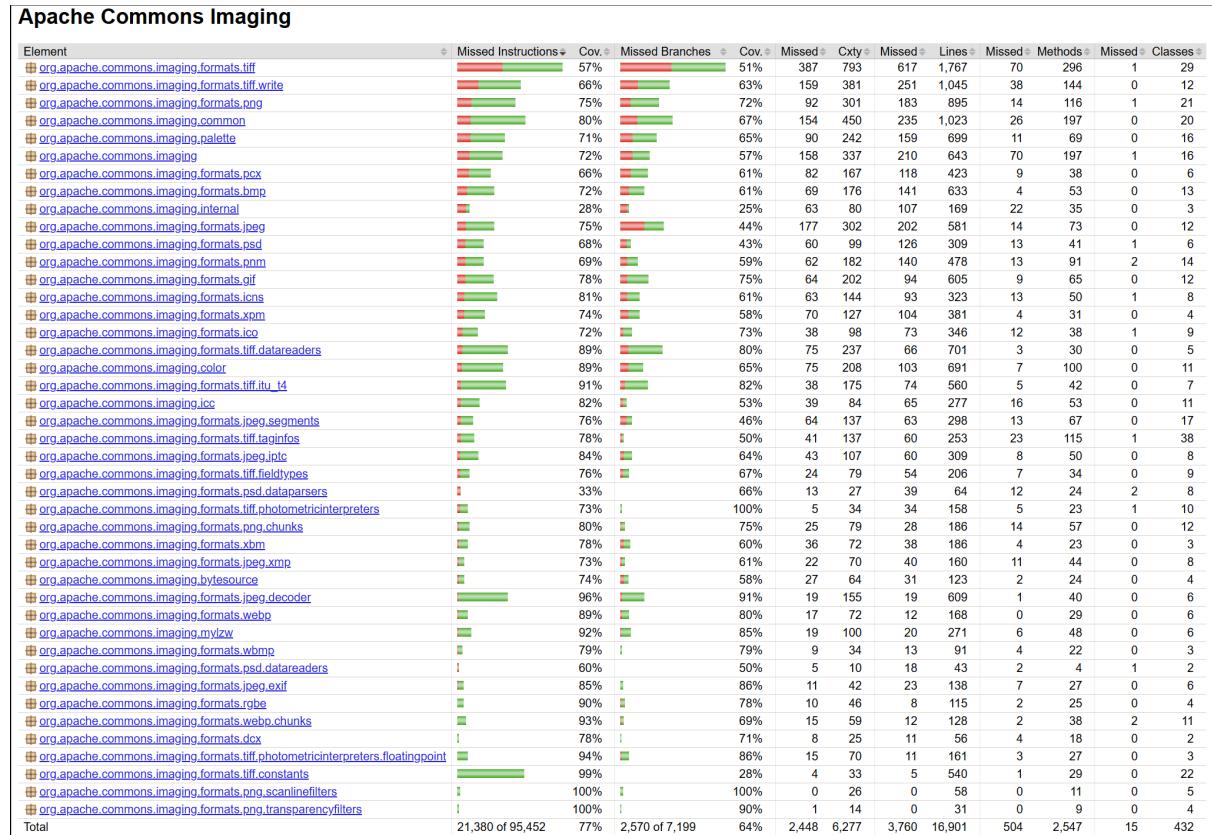


Figure 5.1: Visualization of Initial Code Coverage Results.

5.3 Final Coverage Results After Test Improvements

After generating additional tests using automated tools and adding manual test cases (details are in chapters 6 & 8), the code coverage improved. The final coverage results are summarized below:

Element	Missed Instructions	I. Cov. (%)	Missed Branches	B. Cov. (%)
Total	18,077 of 95,452	81%	2,327 of 7,199	67%

Table 5.2: Final Code Coverage Results After Test Improvements.

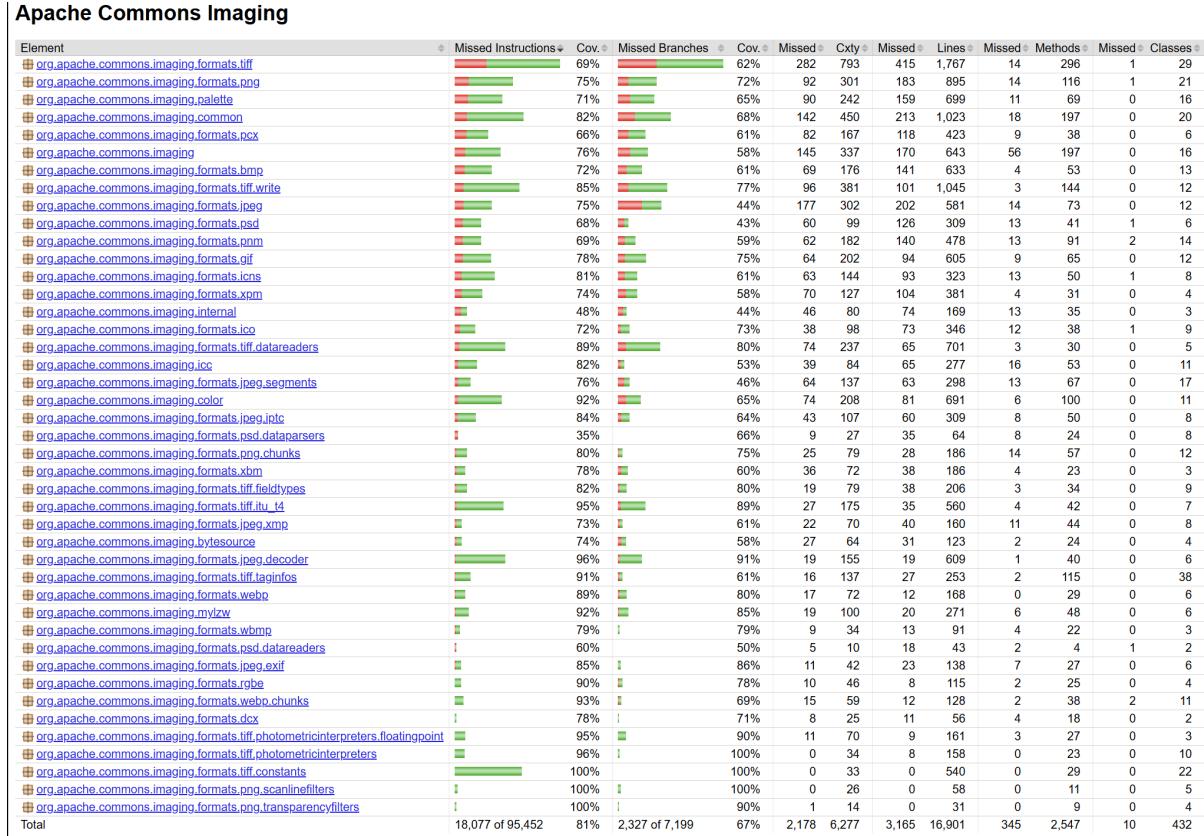


Figure 5.2: Visualization of Final Code Coverage Results.

The results demonstrate some improvements in test coverage and a reduction in missed methods and classes, indicating better validation of the codebase.

6 Mutation Testing with PiTest

Mutation testing is performed to evaluate the effectiveness of the test cases by introducing small modifications (mutations) to the code and ensuring that the tests can detect the changes.

After adding the necessary plugins and dependencies, the following command was run to initiate the mutation testing process:

Listing 6.1: Mutation Testing Command

```
mvn org.pitest:pitest-maven:mutationCoverage
```

6.1 Mutation Test Results

The overall mutation test results for the commons-imaging project are as follows:

Metric	Classes	Line Cov.	Mutation Cov.	Survived Mut.	Test Strength
Overall	69	76% (2973/3904)	56% (1993/3564)	44% (1571/3564)	70% (1993/2839)

Table 6.1: Mutation Test Results Overview.

The number of survived mutants for the whole project is calculated as:

$$\text{Survived Mutants} = \text{Total Mutants} - \text{Killed Mutants}$$

$$\text{Survived Mutants} = 3564 - 1993 = 1571$$

Pit Test Coverage Report

Project Summary

Number of Classes	Line Coverage	Mutation Coverage	Test Strength
69	76% 2973/3904	56% 1993/3564	70% 1993/2839

Breakdown by Package

Name	Number of Classes	Line Coverage	Mutation Coverage	Test Strength
org.apache.commons.imaging	10	68% 432/636	43% 154/360	67% 154/229
org.apache.commons.imaging.bytesource	2	74% 93/125	64% 68/107	82% 68/83
org.apache.commons.imaging.color	11	85% 588/693	54% 443/819	64% 443/695
org.apache.commons.imaging.common	18	77% 791/1032	66% 787/1185	77% 787/1022
org.apache.commons.imaging.jcc	6	80% 218/273	10% 12/117	15% 12/81
org.apache.commons.imaging.internal	3	35% 62/175	19% 21/113	49% 21/43
org.apache.commons.imaging.mylzw	5	93% 253/273	76% 169/223	82% 169/206
org.apache.commons.imaging.palette	14	77% 536/697	53% 339/640	71% 339/480

Report generated by [PIT 1.15.2](#)

Enhanced functionality available at [arcmutate.com](#)

Figure 6.1: Visualization of PIT Coverage Results.

6.2 Detailed Analysis: ColorTools.java

The results for `ColorTools.java` before and after adding additional tests (`ColorToolsTest.java`) were as follows:

Metric	Classes	Line Cov.	Mutation Cov.	Survived Mut.	Test Strength
Before	1	0% (0/62)	0% (0/29)	100% (29/29)	Undefined (0/0)
After	1	60% (37/62)	62% (18/29)	38% (11/29)	90% (18/20)

Table 6.2: Mutation Test Results for `ColorTools.java`.

The following images show the progress made in terms of test coverage for `ColorTools.java`:



Figure 6.2: Initial Coverage and Mutation Results for `ColorTools.java`.



Figure 6.3: Improved Coverage and Mutation Results for `ColorTools.java`.

While 100% line coverage wasn't achieved, this improvement shows a significant increase in coverage, reducing the number of surviving mutants and improving overall test strength.

7 Performance Testing with JMH

Performance tests are implemented using JMH to identify performance bottlenecks in the commons-imaging project.

7.1 Benchmark Setup

The directory containing the benchmarks is located at:

```
.../src/test/java/org/apache/commons/imaging/benchmark
```

The following files are part of the benchmark setup:

- `ImagingBenchmark.java`
- `sample.jpg`
- `sample_baseline.jpg`
- `sample.png`
- `sample.tiff`
- `sample.bmp`

The benchmarks were run by first executing the following Maven command:

```
mvn clean install
```

Then, the `ImagingBenchmark` class was run with the following configuration:

Listing 7.1: ImagingBenchmark-configuration

```
@BenchmarkMode(Mode.AverageTime)
@OutputTimeUnit(TimeUnit.MILLISECONDS)
@State(Scope.Thread)
@Fork(1)
@Warmup(iterations = 3)
@Measurement(iterations = 5)
```

7.2 Stress Testing Results

The most resource-intensive components were stress-tested using JMH, and the results are summarized in the table below:

Benchmark	Mode	Score	Units
ImagingBenchmark.testBmpDecoding	avgt	1.157	ms/op
ImagingBenchmark.testBmpImageInfoRetrieval	avgt	0.300	ms/op
ImagingBenchmark.testBmpImageWriting	avgt	8.348	ms/op
ImagingBenchmark.testImageInfoRetrieval	avgt	0.476	ms/op
ImagingBenchmark.testImageResizing	avgt	36.858	ms/op
ImagingBenchmark.testJpegDecoding	avgt	212.311	ms/op
ImagingBenchmark.testPngEncoding	avgt	132.896	ms/op
ImagingBenchmark.testTiffMetadataExtraction	avgt	1.217	ms/op

Table 7.1: Stress Testing Results of Commons-Imaging Project.

Benchmark	Mode	Cnt	Score	Error	Units
ImagingBenchmark.testBmpDecoding	avgt		1.157		ms/op
ImagingBenchmark.testBmpImageInfoRetrieval	avgt		0.300		ms/op
ImagingBenchmark.testBmpImageWriting	avgt		8.348		ms/op
ImagingBenchmark.testImageInfoRetrieval	avgt		0.476		ms/op
ImagingBenchmark.testImageResizing	avgt		36.858		ms/op
ImagingBenchmark.testJpegDecoding	avgt		212.311		ms/op
ImagingBenchmark.testPngEncoding	avgt		132.896		ms/op
ImagingBenchmark.testTiffMetadataExtraction	avgt		1.217		ms/op

Process finished with exit code 0

Figure 7.1: Benchmark Results of Selected Methods of Commons-Imaging Project.

8 Automated Test Generation

Automated tests were generated using Randoop to improve coverage on poorly tested components. These tests focus on edge cases and scenarios not previously covered.

8.1 Low Coverage Packages

From the JaCoCo report, we identified that three packages had low coverage:

- internal
- formats/psd/dataparsers
- formats/tiff

The report also highlights the classes within these packages that need additional test coverage, as shown in the image below:

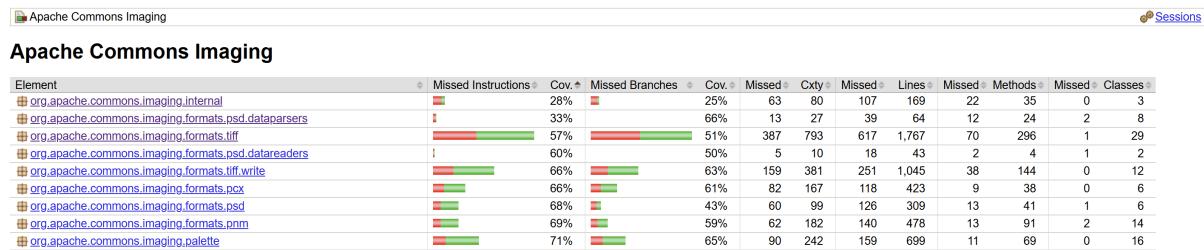


Figure 8.1: JaCoCo report showing low coverage in specific packages.

8.2 Test Generation Process

To generate tests for all the classes within these packages, we first listed the classes using the following commands:

```
cd target/classes
find org/apache/commons/imaging/internal -name "*.class" | sed 's/\//./g' | sed
    ↪ 's/\.class$//'
find org/apache/commons/imaging/formats/psd/dataparsers -name "*.class" | sed
    ↪ 's/\//./g' | sed 's/\.class$//'
find org/apache/commons/imaging/formats/tiff -name "*.class" | sed 's/\//./g' | sed
    ↪ 's/\.class$//'
```

These commands search for all ‘.class’ files in the specified directories and convert them into fully qualified class names, which are then copied and saved to a text file called `Randoop_Test_Classes.txt` manually.

8.3 Running Randoop

Once the class names were added to the `Randoop_Test_Classes.txt` file, we ran the following command to generate the tests:

```
java -cp "randoop-all-4.3.2.jar;target/classes" randoop.main.Main gentests
    ↪ $(Get-Content Randoop_Test_Classes.txt | ForEach-Object \{ "--testclass=$\_\\"})
    ↪ ) --time-limit=200
    ↪ --junit-output-dir=src/test/java/org/apache/commons/imaging/randoop/
```

This command uses the Randoop tool to generate tests for the classes listed in the text file. The generated tests are stored in the `randoop` folder and are written in JUnit format.

8.4 Test Coverage Improvement

After generating the tests, we observed an improvement in coverage, as shown in the following image:

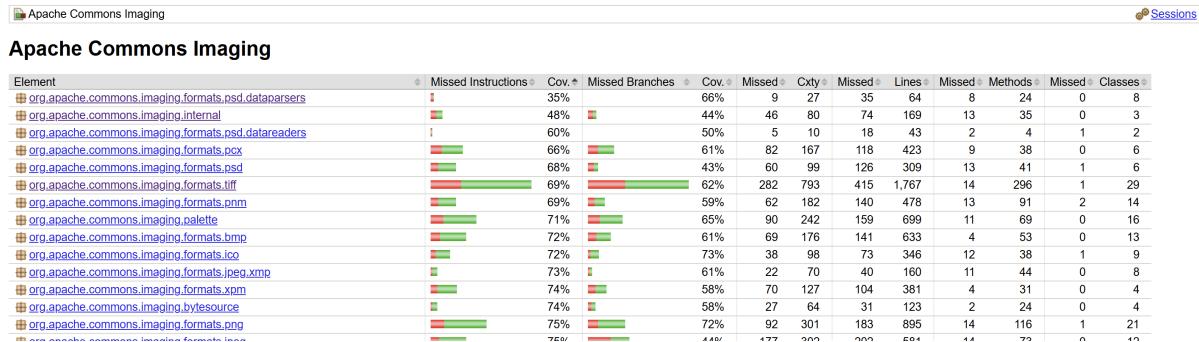


Figure 8.2: Coverage improvement after generating tests using Randoop.

Metric	Instruction Coverage	Branch Coverage
Before	28%	25%
After	48%	44%

Table 8.1: Coverage for `org.apache.commons.imaging.internal`.

Metric	Instruction Coverage	Branch Coverage
Before	33%	66%
After	35%	66%

Table 8.2: Coverage for `org.apache.commons.imaging.formats.psd.dataparsers`.

Metric	Instruction Coverage	Branch Coverage
Before	57%	51%
After	69%	62%

Table 8.3: Coverage for `org.apache.commons.imaging.formats.tiff`.

8.5 Challenges and Limitations

While this method improved coverage, it has some limitations:

- The generated tests often contain redundancy.
- It does not guarantee sufficient coverage for all methods.
- The process generates more tests than are necessary, leading to inefficiency.

Despite these limitations, automated test generation with Randoop provides a valuable first step in improving test coverage, especially for poorly tested components.

9 Security Analysis with Snyk

9.1 Steps to Run Snyk

To perform security analysis using Snyk, follow these steps:

```
npm install -g snyk
snyk auth
snyk test
snyk test --json > snyk-report.json
```

9.2 Vulnerability Report

There were no vulnerabilities found for the commons-imaging project, as seen in the image below. However, there are some vulnerable paths in the web app backend Dockerfile, which will be reported instead.

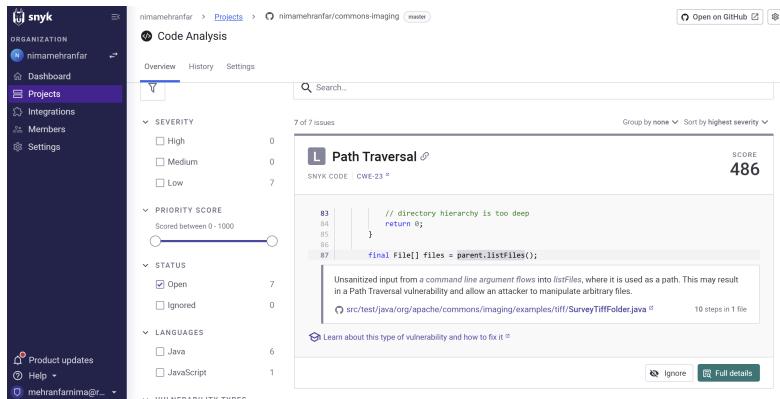


Figure 9.1: Snyk report showing no vulnerabilities for commons-imaging.

By changing the OpenJDK version to a newer one, many of the vulnerabilities were resolved, as shown in the image below:

Recommendations for upgrading the base image			
	BASE IMAGE	VULNERABILITIES	SEVERITY
Current image	openjdk:21-jdk	87	0 [C] 35 [H] 42 [M] 10 [L]
Major upgrades	openjdk:24-jdk	35	0 [C] 6 [H] 29 [M] 0 [L]
View docs			Open a fix PR

Figure 9.2: Vulnerabilities reduced after updating the OpenJDK version.

9.3 Security Issues Found

During the analysis, the following security issues were identified:

- 7 "low" level issues, all of which were related to unsanitized input.
- These issues seem to be false positives because the project is an API and utility, and there is no allowed directory. Users can use any directory they want.

- However, if we want to limit the user to their current directory, we can fix these issues. For example, unsanitized input from a command-line argument was found in the file `src/test/java/org/apache/commons/imaging/examples/tiff/SurveyTiffFolder.java:collectPaths.class`, leading to a potential Path Traversal vulnerability.

Below is an image showing the identified vulnerability:

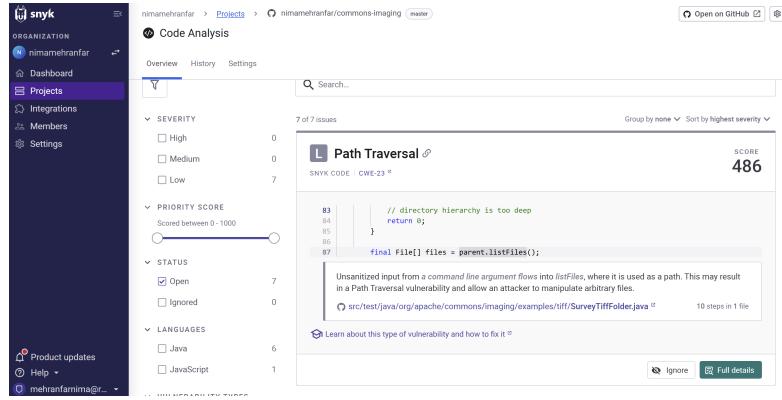


Figure 9.3: Unsanitized input vulnerability found in SurveyTiffFolder.java:collectPaths.

To fix this issue, we can ensure that the file paths remain within the base directory, as shown in the code snippet below:

Listing 9.1: SurveyTiffFolder.java:collectPath.class fix

```
private static final String BASE_DIRECTORY = System.getProperty("user.dir");

private static boolean isWithinBaseDirectory(File file) throws IOException {
    // Get the canonical path of the file
    String canonicalPath = file.getCanonicalPath();
    String baseDirectoryPath = new File(BASE_DIRECTORY).getCanonicalPath();

    // Check if the canonical path starts with the base directory path
    return canonicalPath.startsWith(baseDirectoryPath);
}

private static int collectPaths() {
    ...
    try {
        // Resolve the file's canonical path to get the absolute path
        File canonicalFile = f.getCanonicalFile();

        // Normalize the path and ensure it is within the base directory
        if (!isWithinBaseDirectory(canonicalFile)) {
            continue; // Skip files outside the base directory
        }

        final String[] temp = Arrays.copyOf(scratch, depth + 1);
        pathList.add(temp);
    } catch (IOException e) {
        // Handle any IOExceptions (e.g., if getCanonicalPath fails)
    }
    ...
}
```

This fix ensures that the file operations are limited to the base directory, preventing path traversal vulnerabilities.

All other identified issues can be fixed in a similar manner if they are not false positives, by sanitizing the inputs and ensuring that file operations are safe.

10 Conclusion

The commons-imaging project underwent various quality assurance practices. Through CI/CD pipeline integration, refactoring based on SonarCloud issues, performance and mutation testing, and security analysis, we have significantly improved the project's code quality and created a test Docker for it. Future work will depend on project developers focusing on addressing any remaining vulnerabilities and expanding the test coverage.

The source code for this project is available at [GitHub Repository](#).

11 References

- SonarCloud Documentation: <https://sonarcloud.io/>.
- Docker Documentation: <https://docs.docker.com/>.
- JaCoCo Documentation: <https://www.jacoco.org/jacoco/>.
- PiTest Documentation: <https://pitest.org/>.
- JMH Documentation: <https://openjdk.java.net/projects/code-tools/jmh/>.
- Randoop Documentation: <http://randoop.github.io/>.
- Snyk Documentation: <https://snyk.io/>.
- ChatGPT: <https://chatgpt.com/>.